

Thread-Level Speculation with Kernel Support

Clemens Hammacher, Kevin Streit,
Andreas Zeller and Sebastian Hack



International Conference on Compiler Construction
March 17th, 2016, Barcelona

Thread-Level Speculation (TLS)

- *unsafe / optimistic* parallel execution
- data or control dependences might exist
- runtime system checks for violations
- rollback & re-execute in case of conflict

Thread-Level Speculation (TLS)

- *unsafe / c*

```
for (auto &task : tasks[today]) {
    task.compute();
    if (!task.success()) {
        fatalError("task failed: " + task.id());
    }
}
```

- data or co

- runtime s

```
for (auto id : taskIds[today]) {
    tasks[id].compute();
}
```

- rollback &

Thread-Level Speculation (TLS)

- *unsafe /*

```
for (auto &task : tasks[today]) {
    task.compute();
    if (!task.success()) {
        fatalError("task failed: " + task.id());
    }
}
```

- data or co

- runtime s

```
for (auto id : taskIds[today]) {
    tasks[id].compute();
}
```

- rollback &

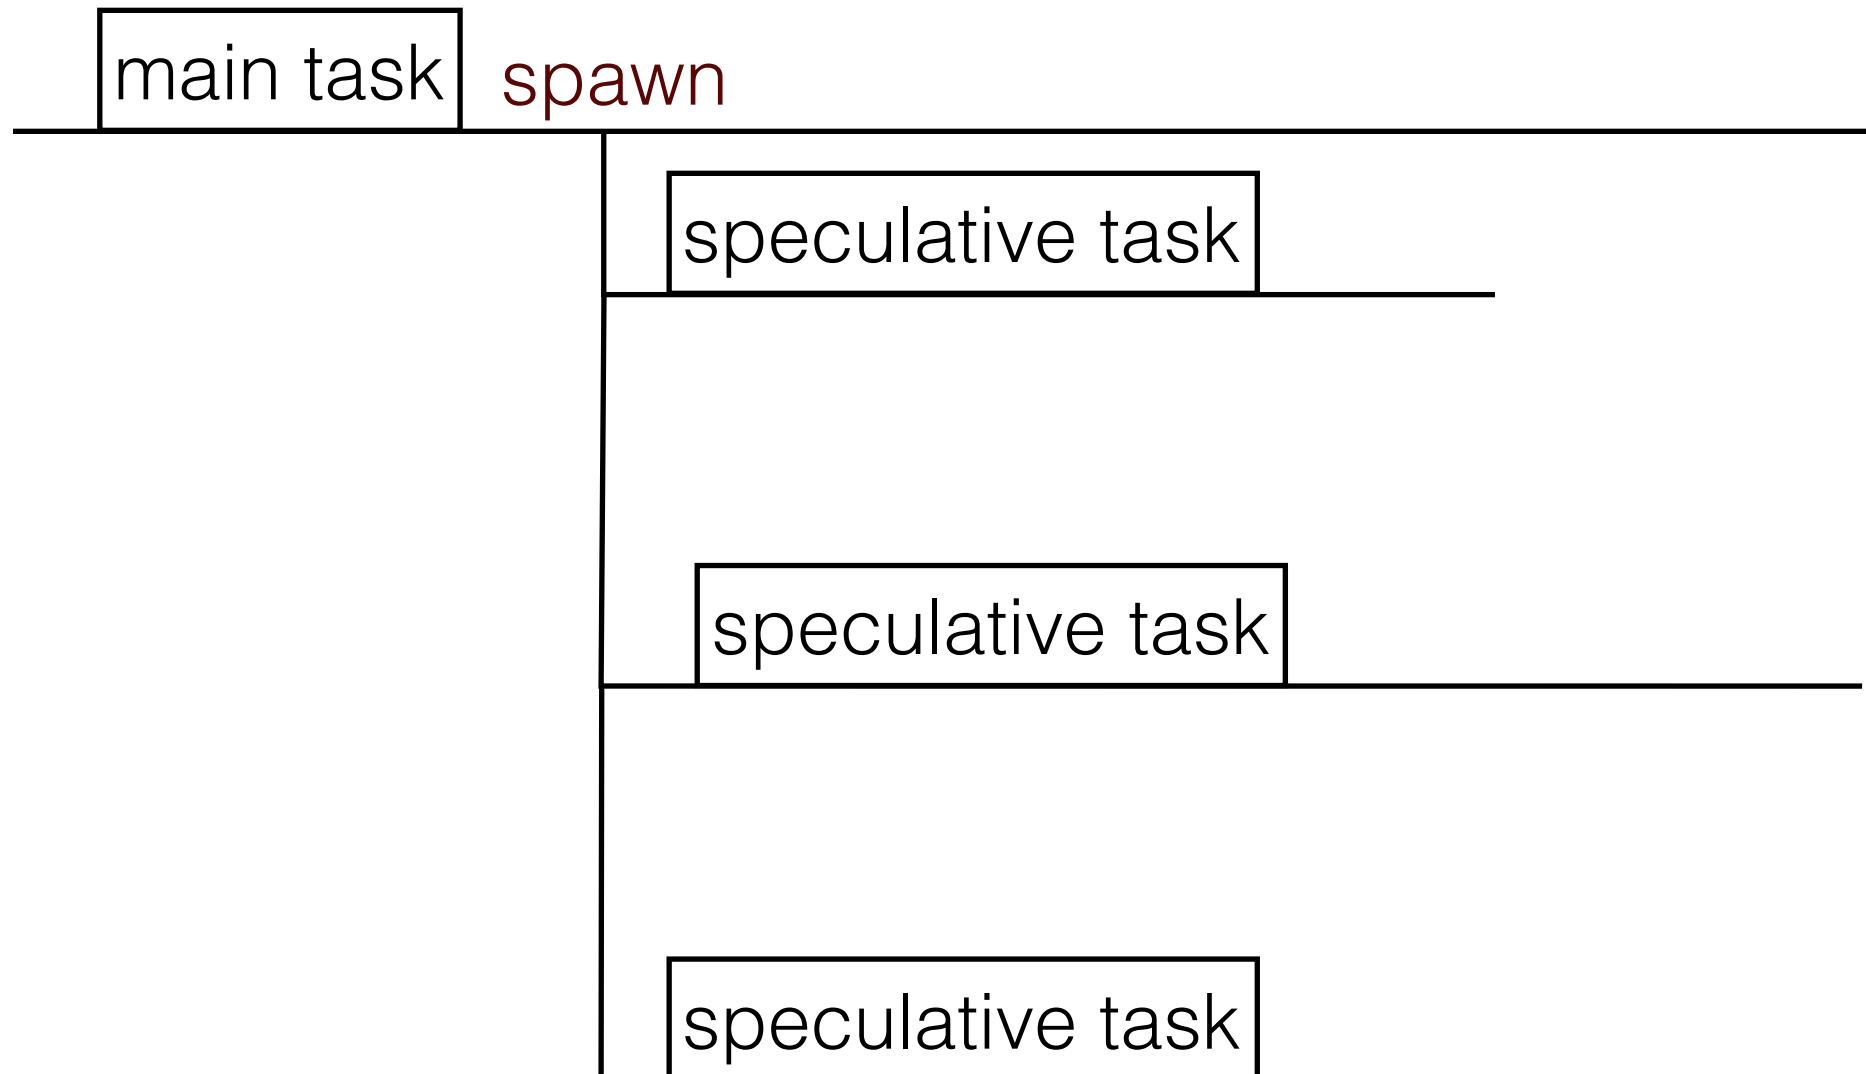
→ manual / guided parallelization

→ automatic parallelization


Execution in TLS

main task

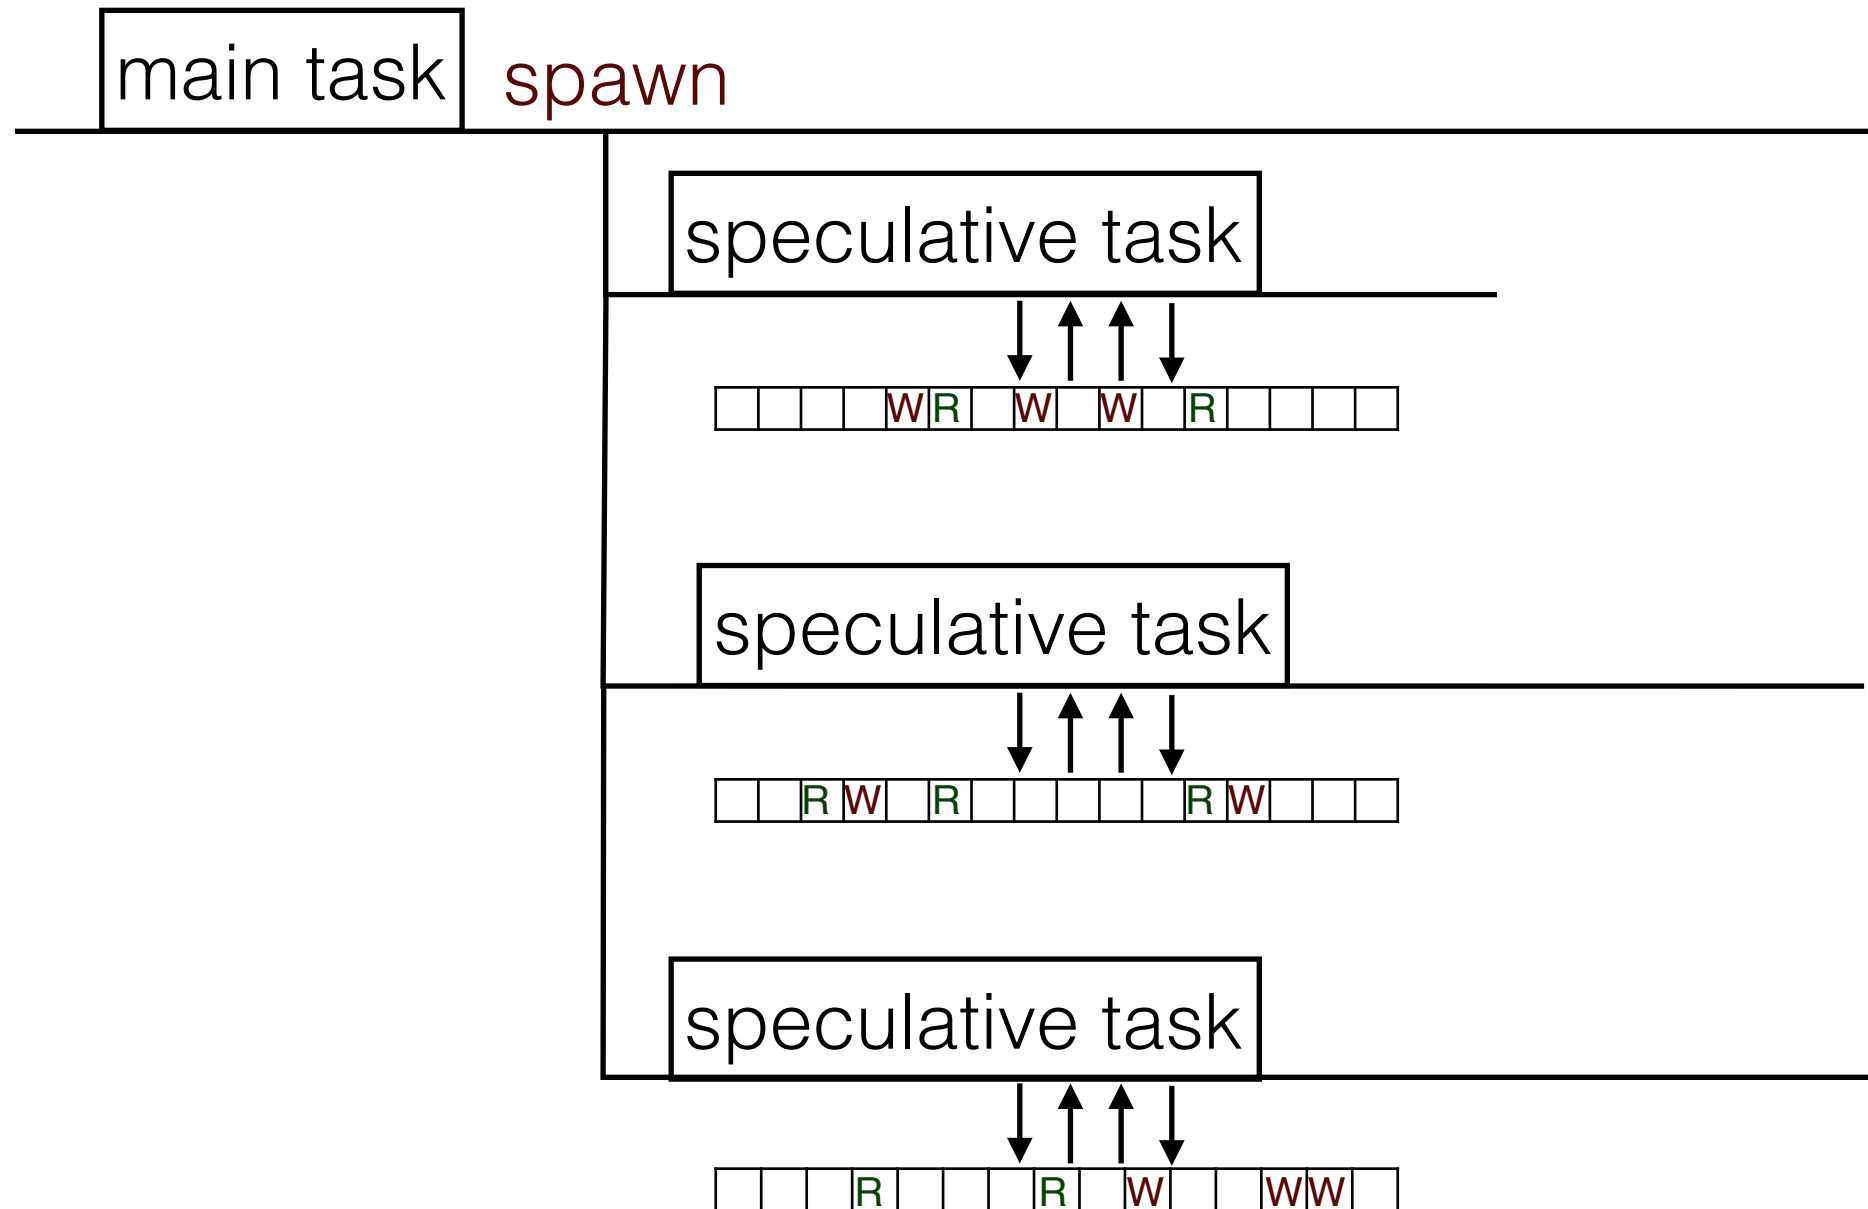
Execution in TLS



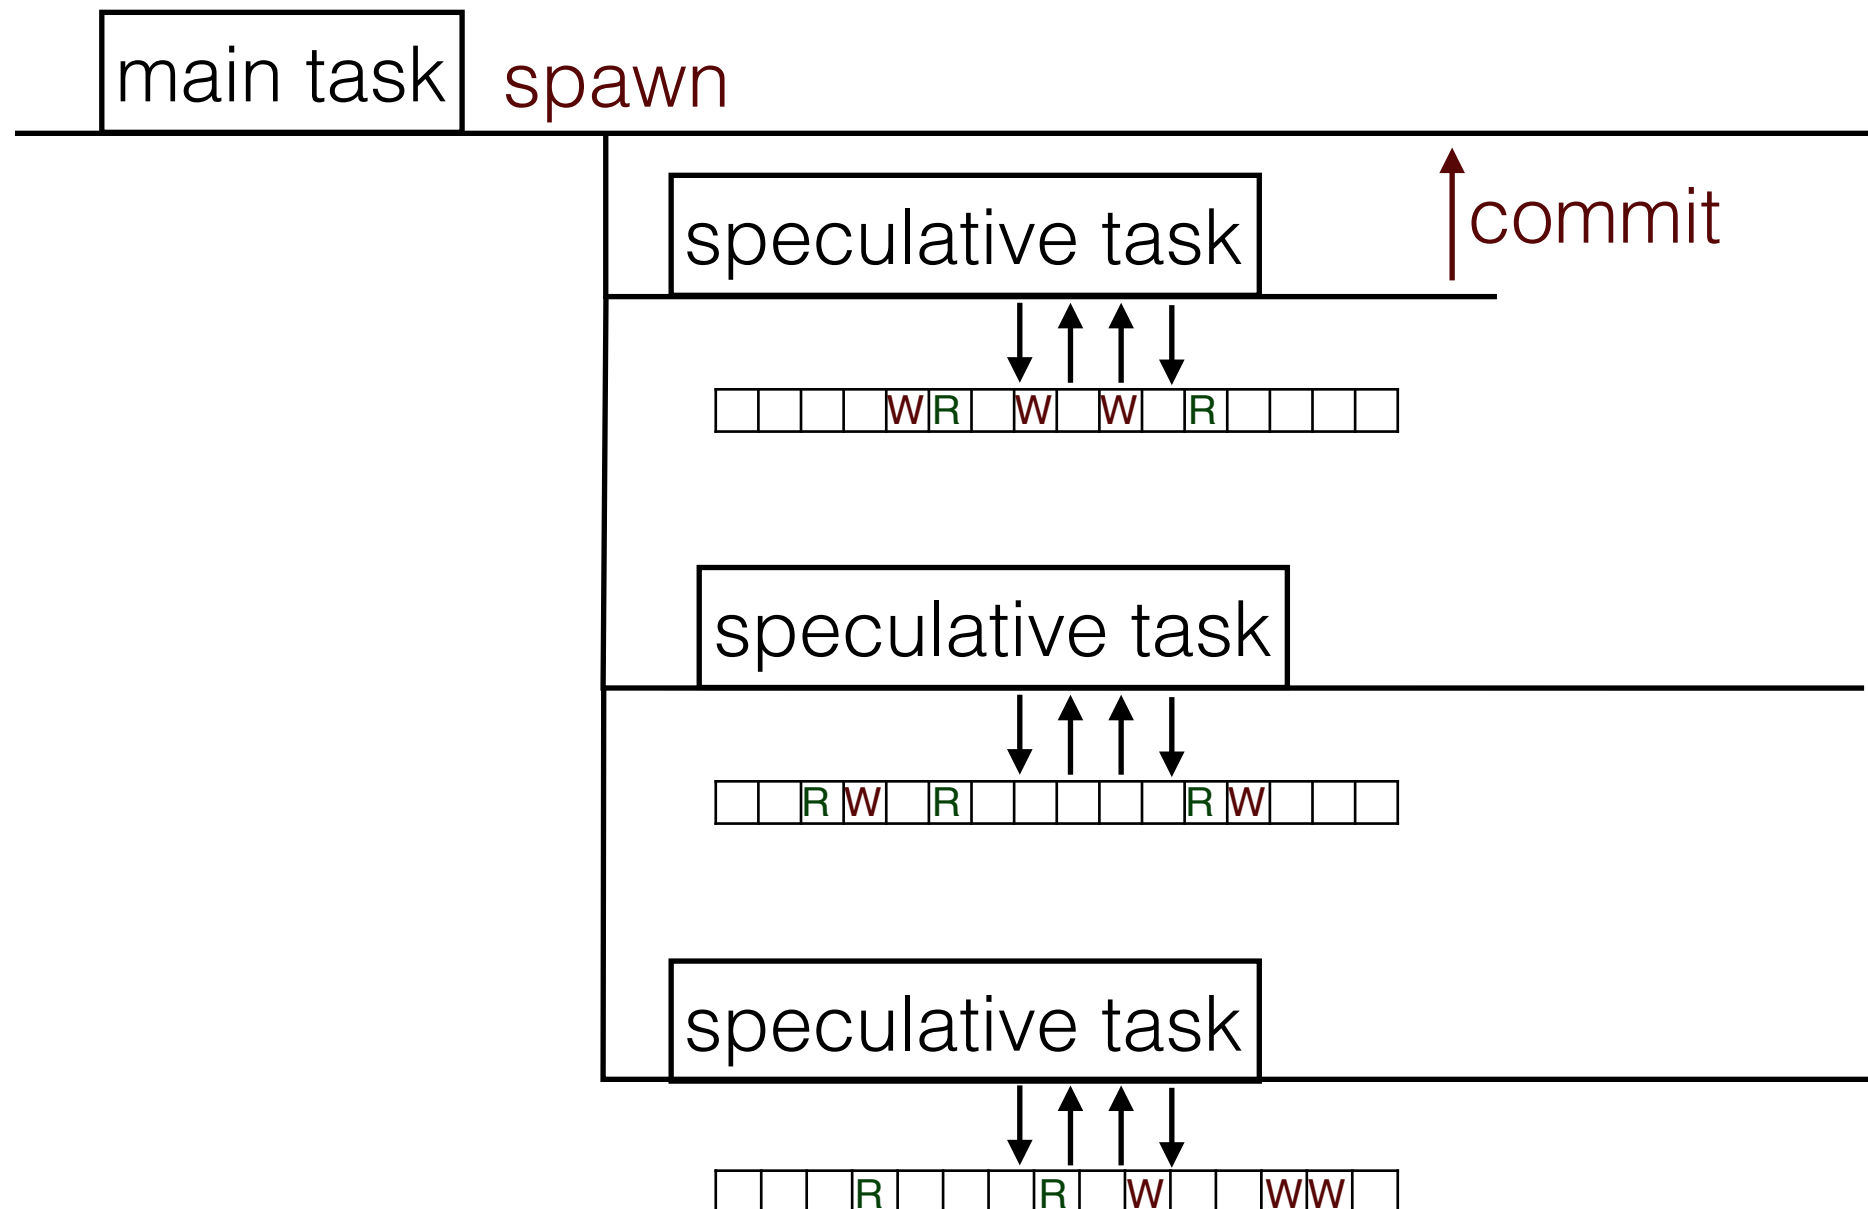
Execution in TLS




Execution in TLS



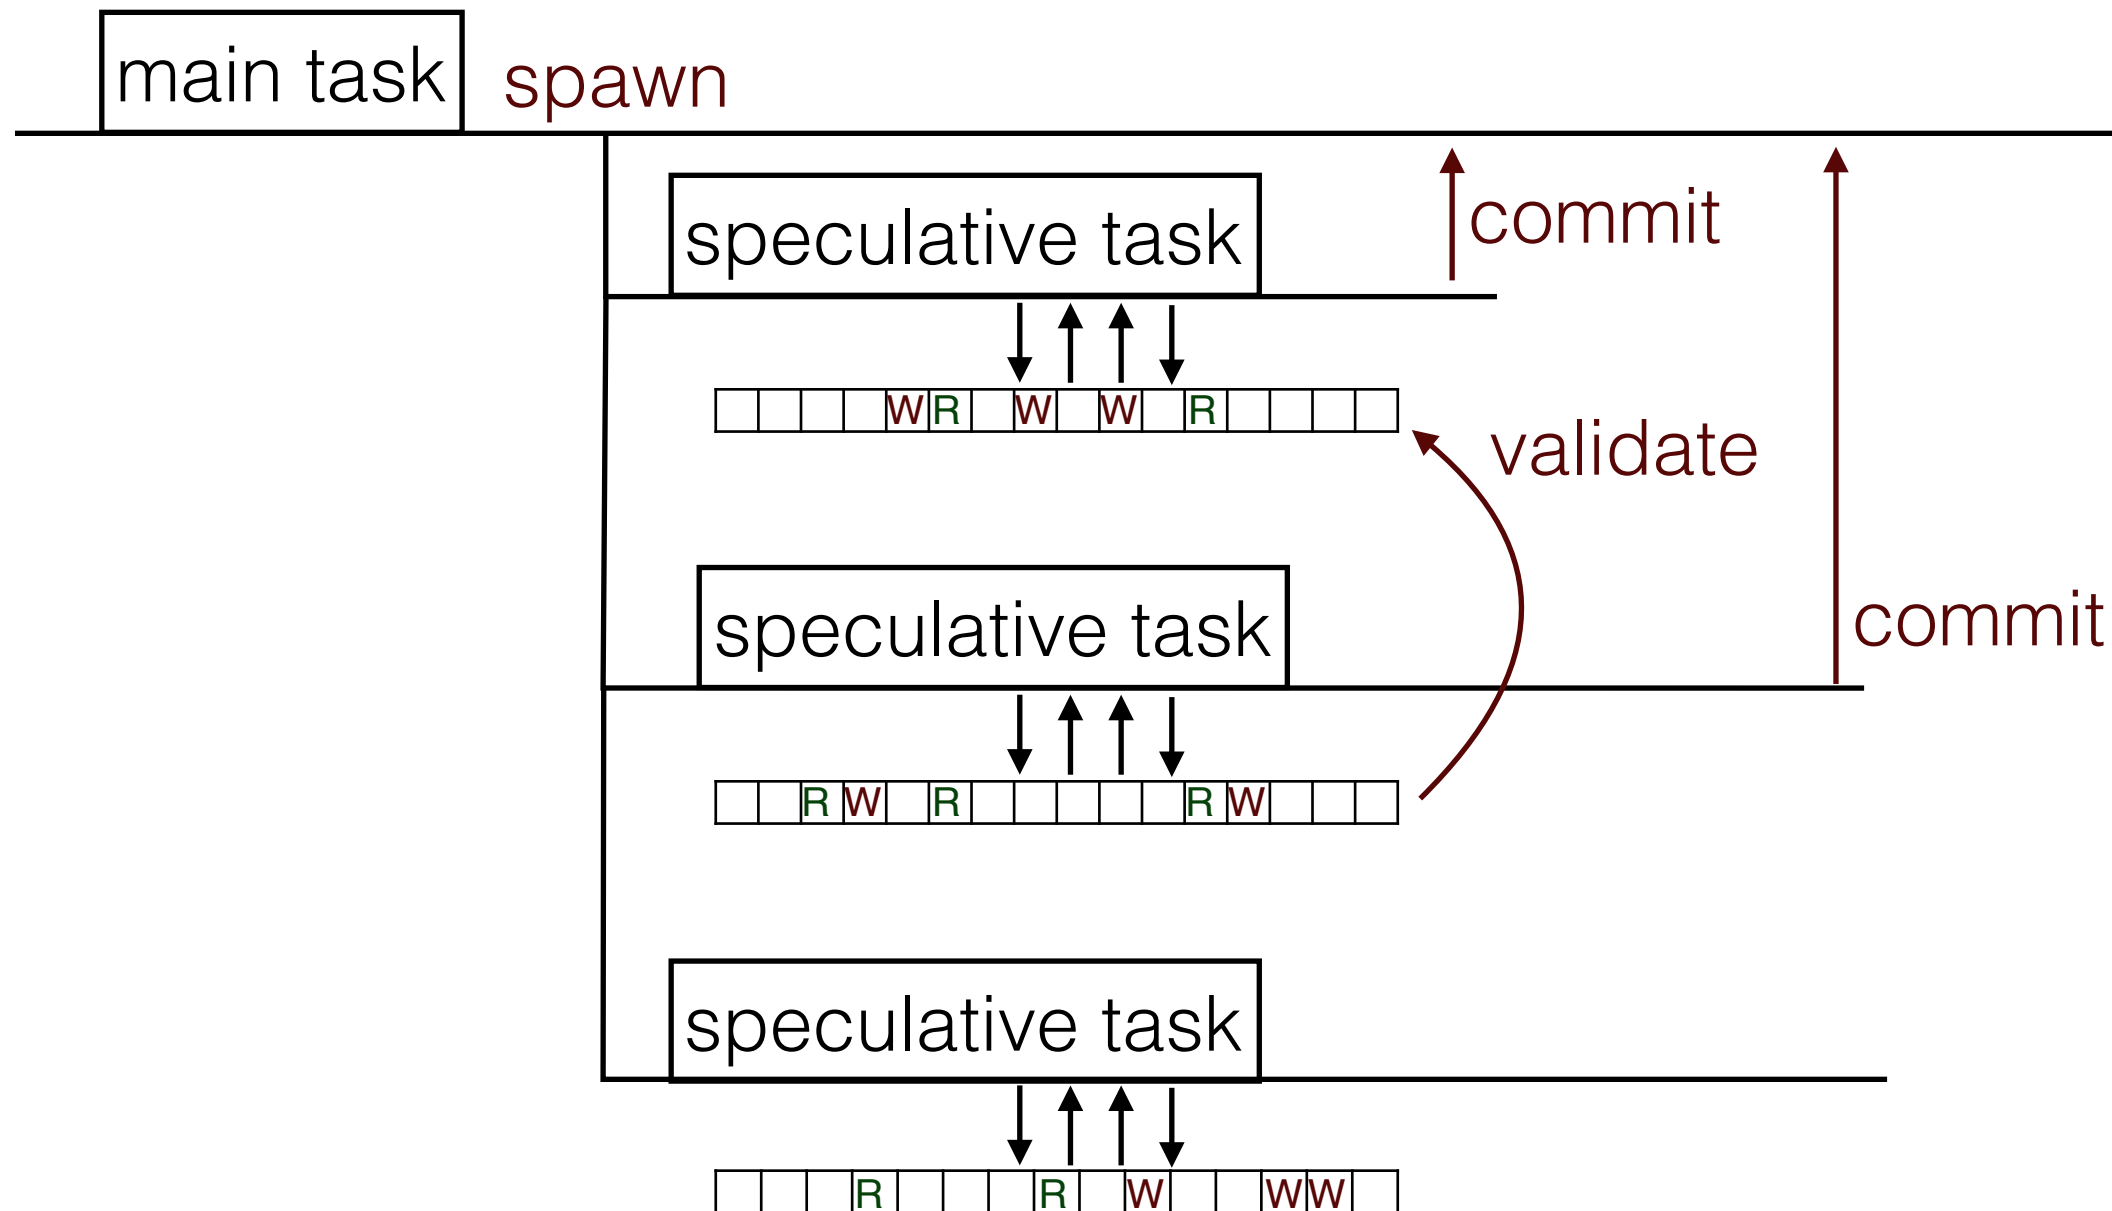
Execution in TLS



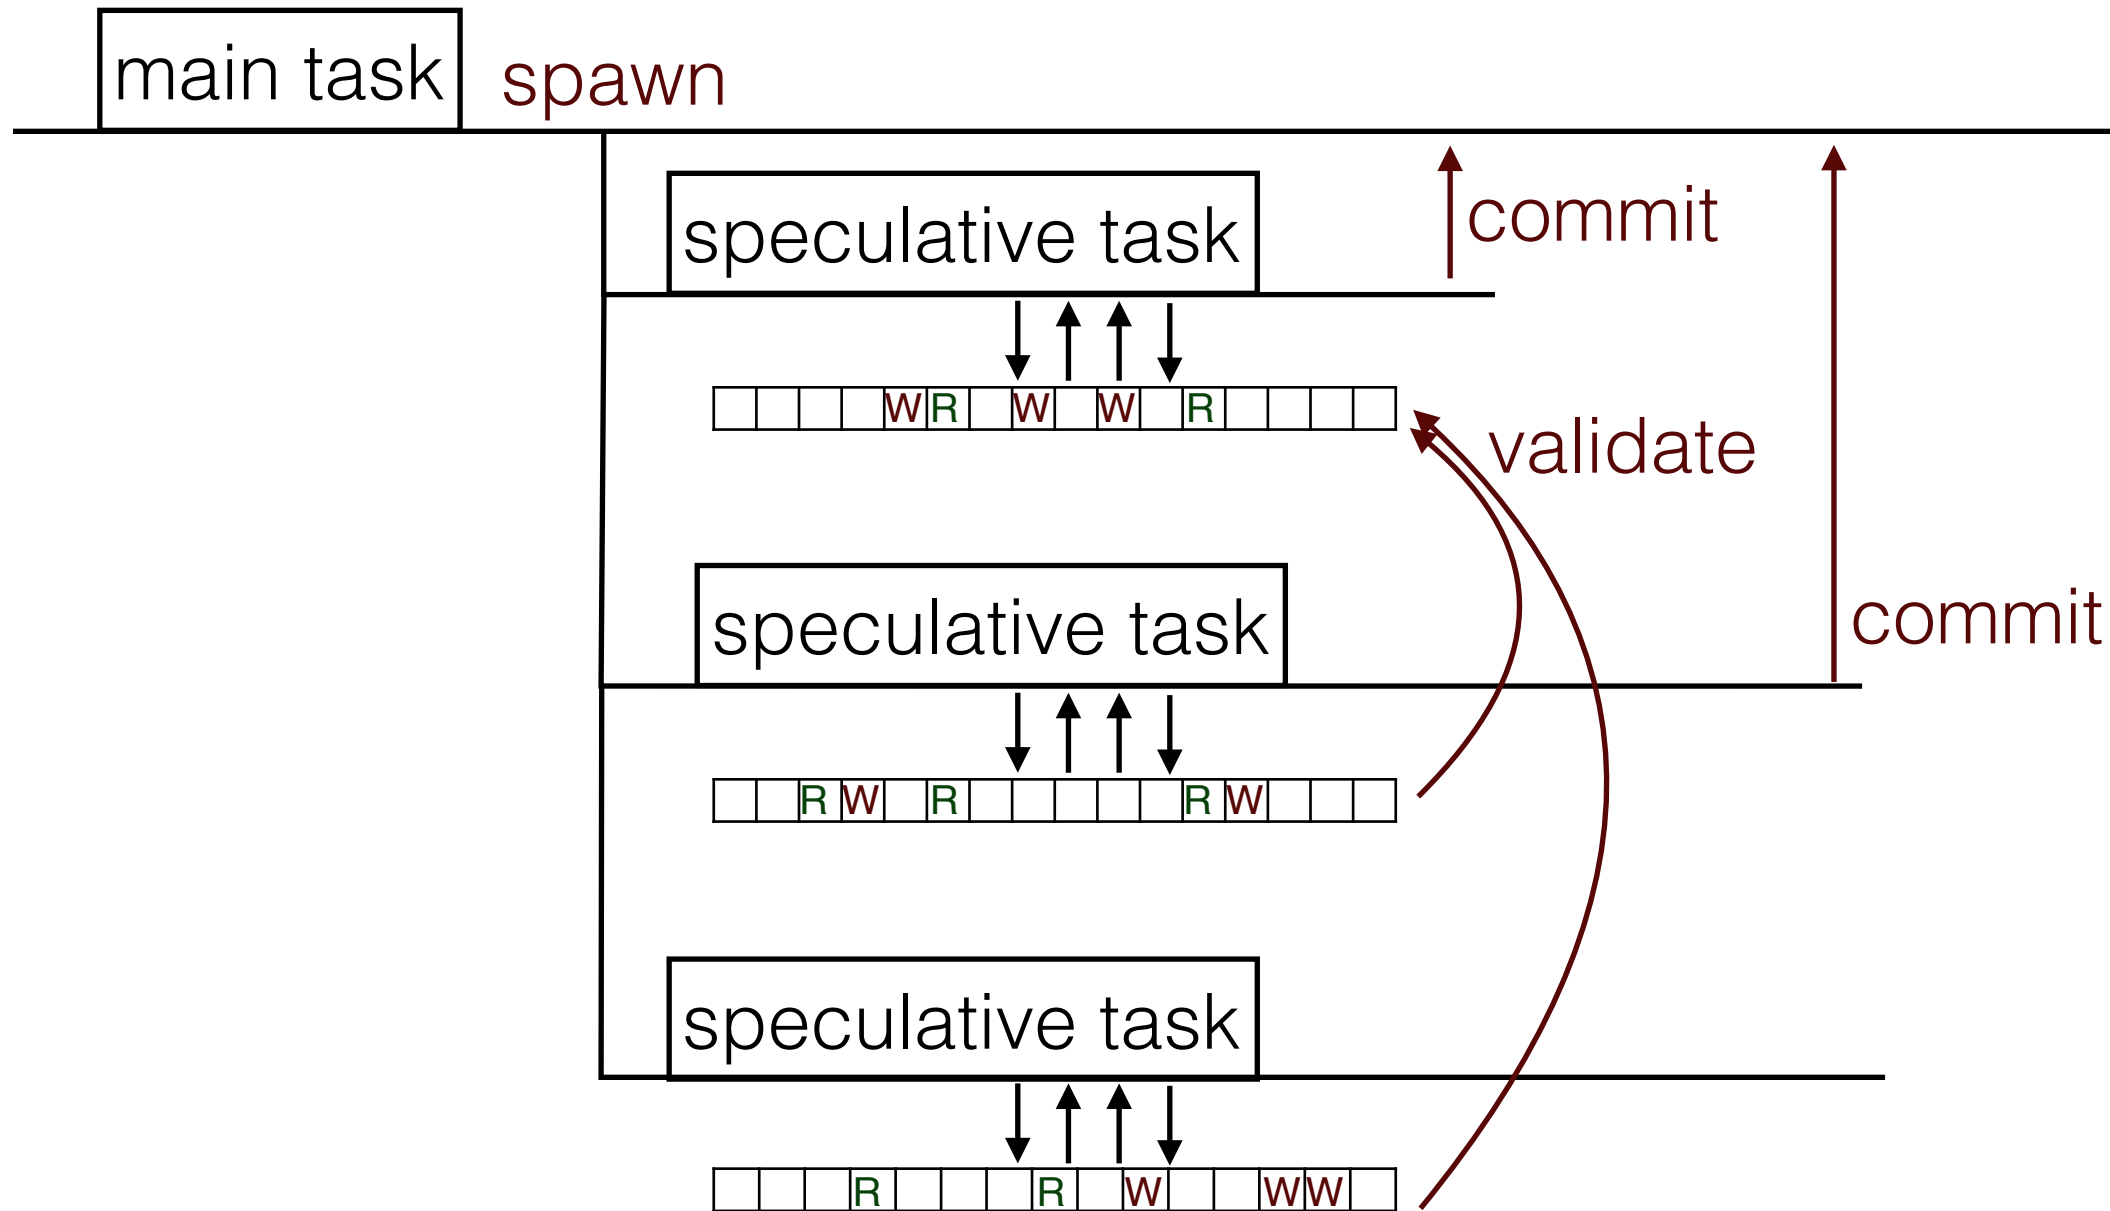
Execution in TLS



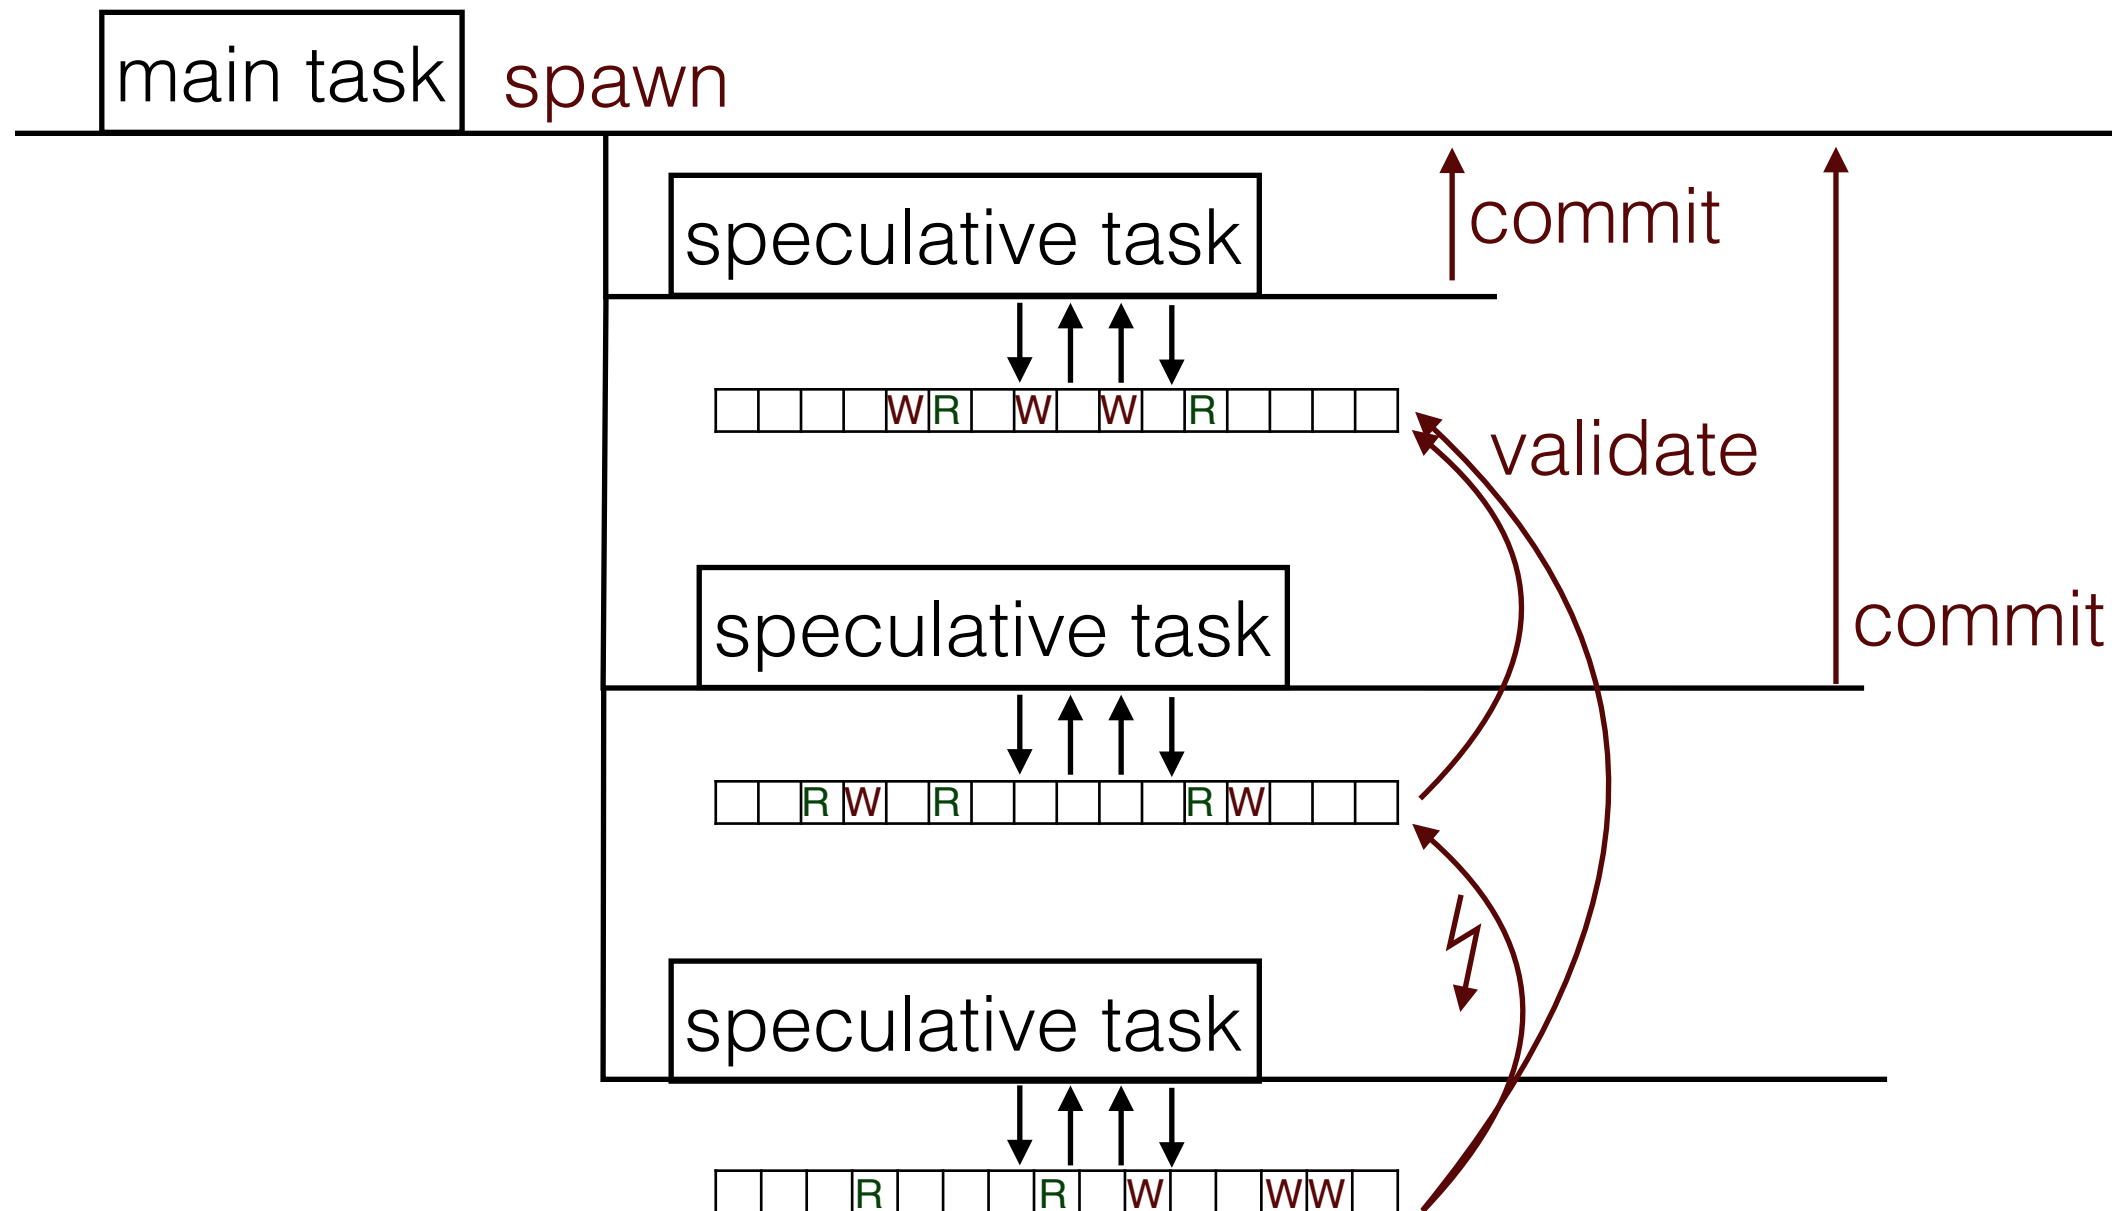
Execution in TLS




Execution in TLS



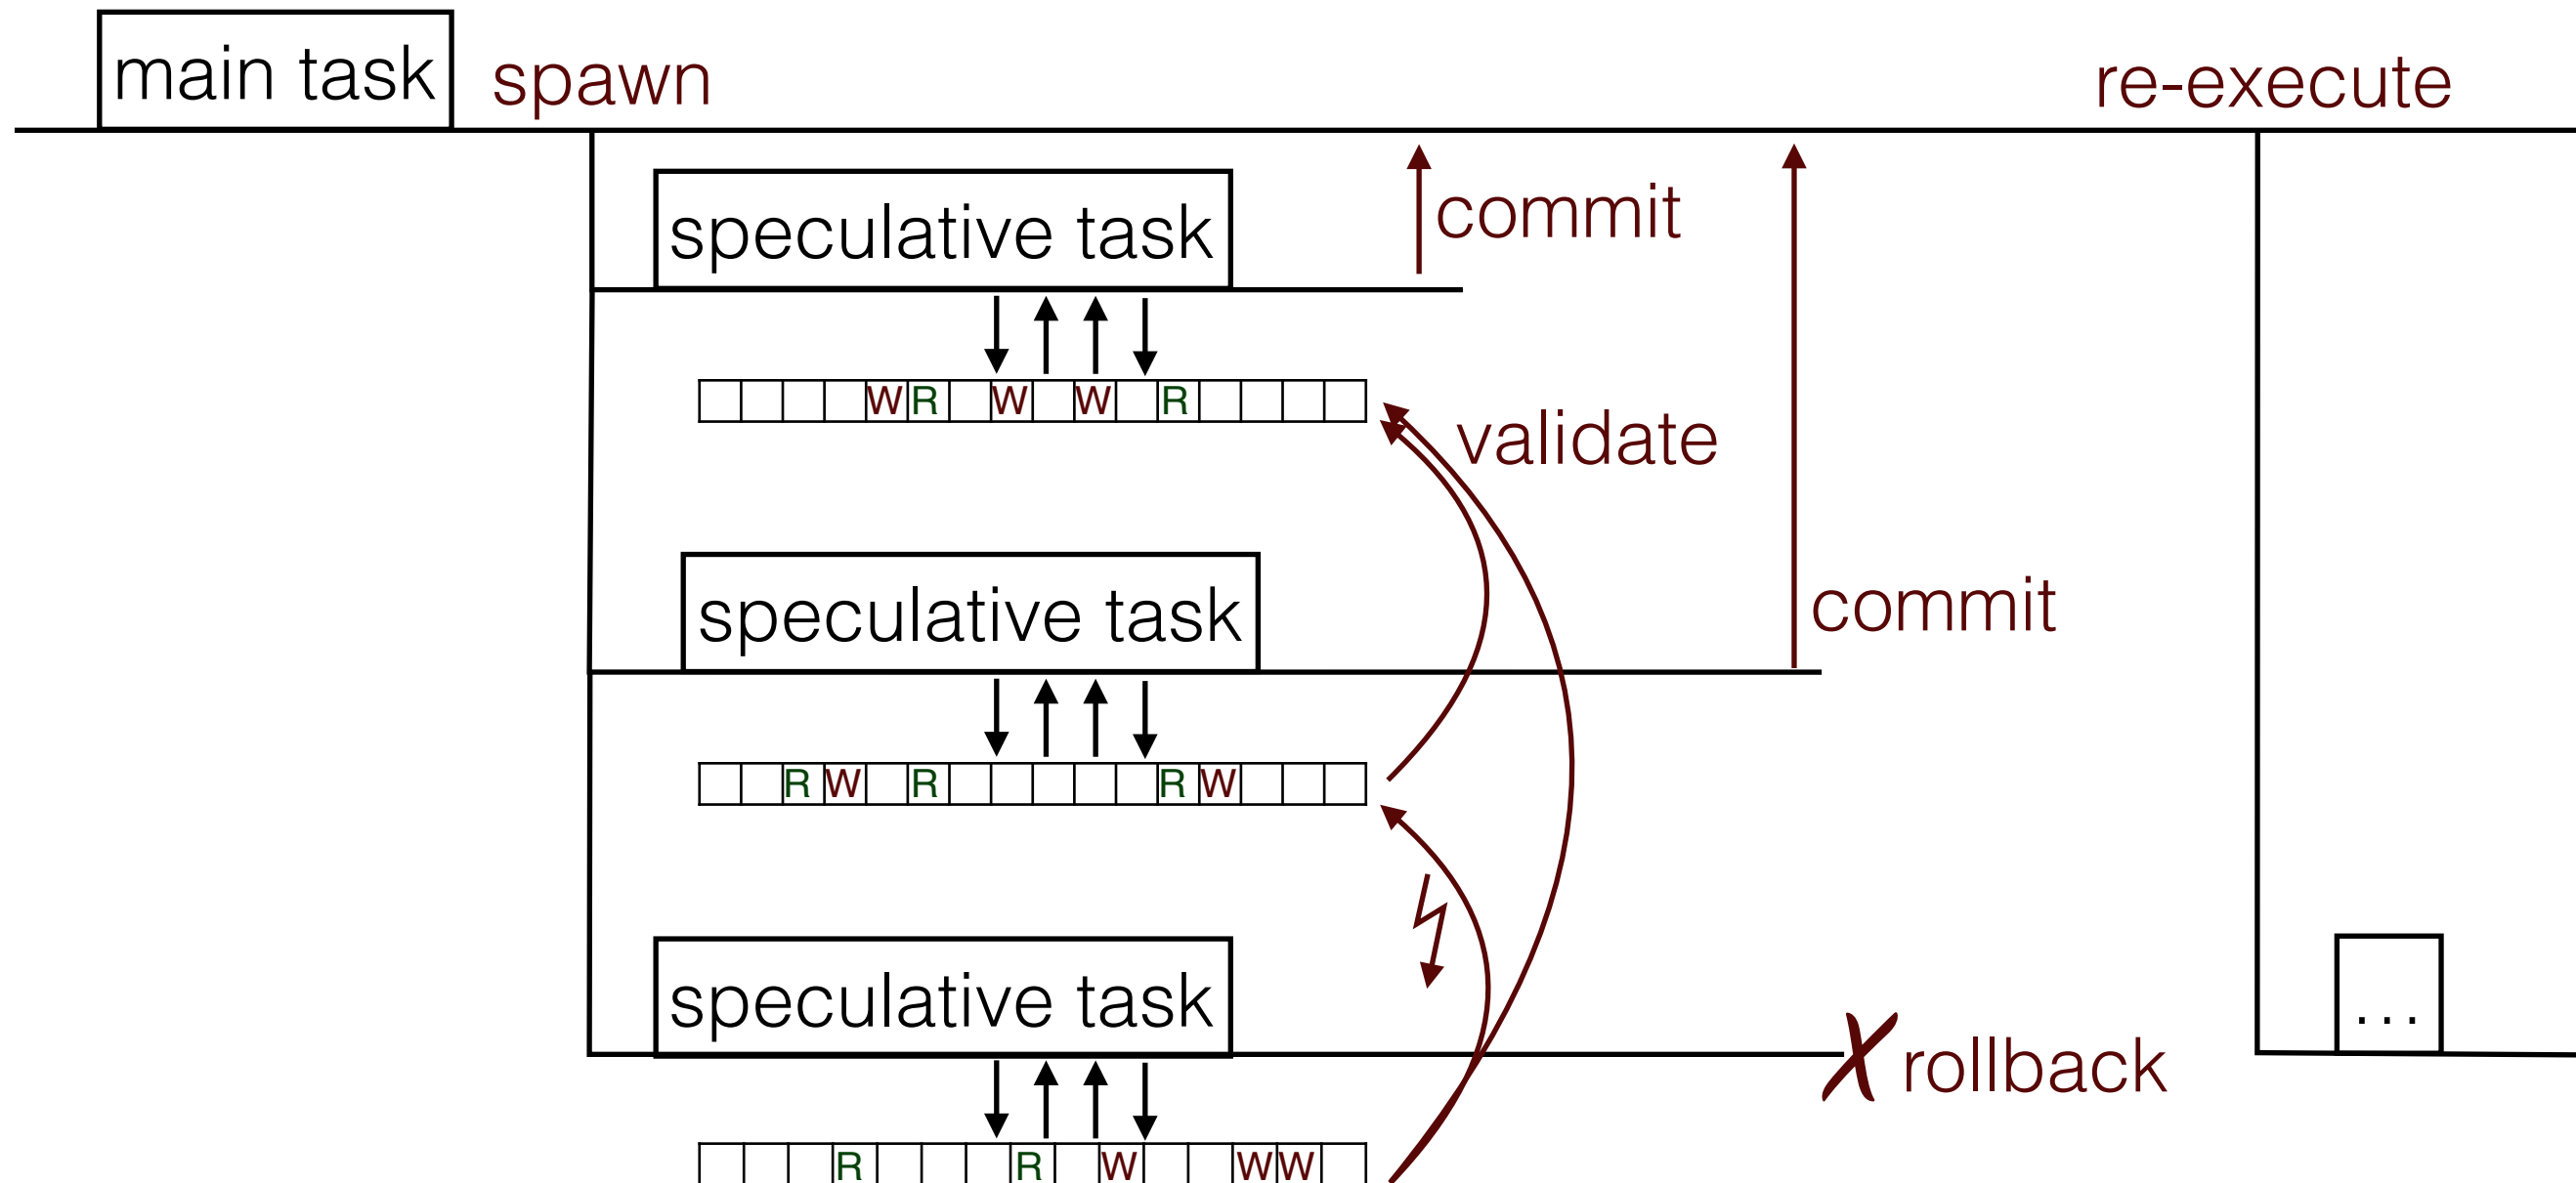
Execution in TLS



Execution in TLS



Execution in TLS



TLS Implementations

S-TLS

VM-TLS

- Thread-Level Speculation in pure software
 - similar to software transactional memory (STM)
 - *explicit* tracking of memory accesses
 - code *instrumentation*
 - *multithreaded* in same address space
 - first described in 1993 by Herlihy and Moss

TLS Implementations

S-TLS

VM-TLS

performance on STM benchmarks
(lots of small transactions)


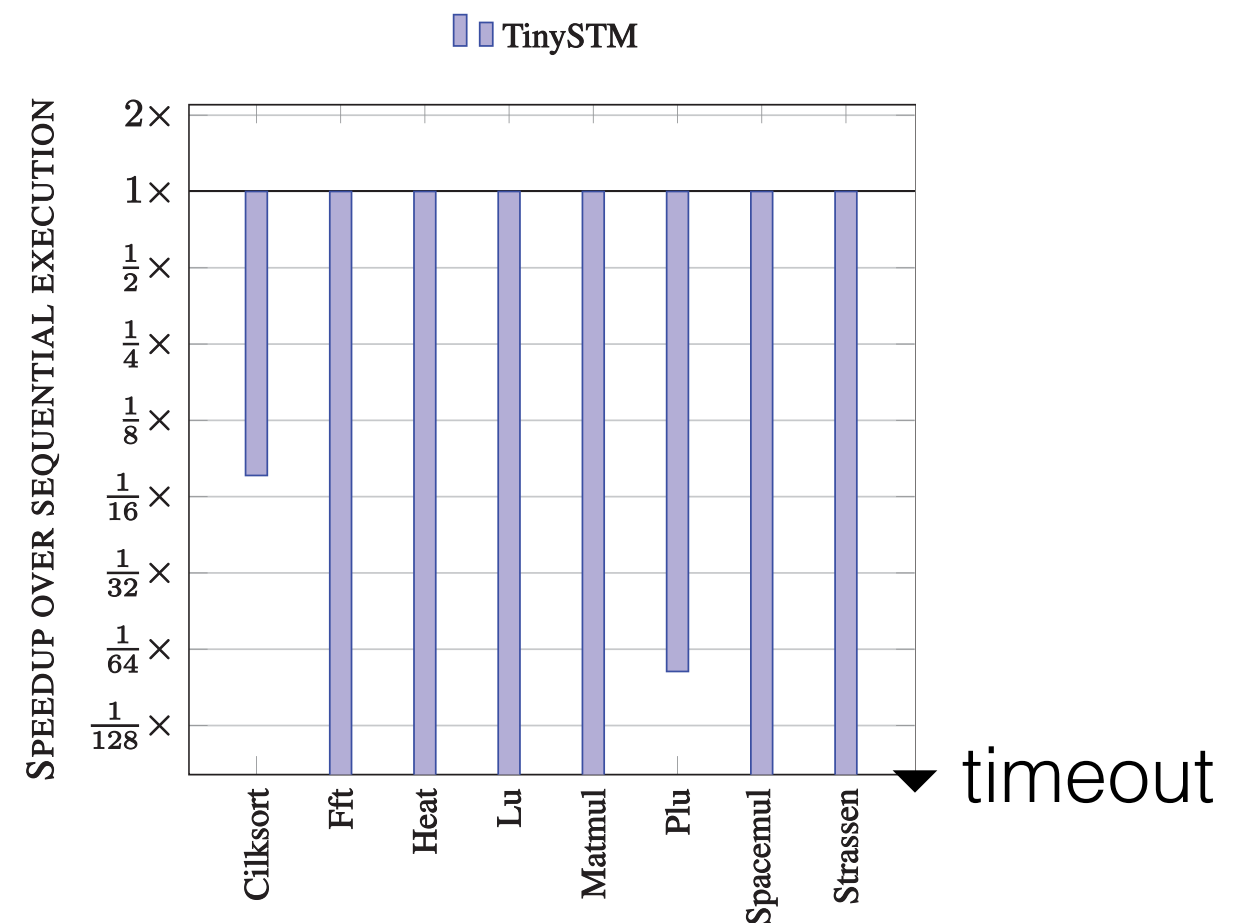


Figure 11. TinySTM performance with 16 core x86

source: Dragojevic et al., 2011, *Communications of the ACM*.

performance on TLS benchmarks
(few large transactions)



TLS Implementations

S-TLS

VM-TLS

- Virtual-memory based approaches
 - *implicit* tracking by protecting memory pages
 - segfault handler logs access & unprotects the page
 - can execute *arbitrary code*, even from compiled libraries
 - *multi-process execution* for separate address spaces
 - introduces in 2001 by Papadimitriou and Mowry

Memory Protection

In Virtual-Memory Based TLS

main process:

VMA 1: r-x

VMA 2: rw-



Memory Protection

In Virtual-Memory Based TLS

main process:

VMA 1: r-x

VMA 2: rw-



speculative process:

VMA 1: r-x


VMA 2: rw-




Memory Protection

In Virtual-Memory Based TLS

main process:




speculative process:




Memory Protection

In Virtual-Memory Based TLS

main process:




speculative process:




Memory Protection

In Virtual-Memory Based TLS

main process:



speculative process:




→ `mprotect(addr, 4096, PROT_READ)`


Memory Protection

In Virtual-Memory Based TLS

main process:



speculative process:




→ `mprotect(addr, 4096, PROT_READ)`


Memory Protection

In Virtual-Memory Based TLS

main process:




speculative process:




Memory Protection

In Virtual-Memory Based TLS

main process:




speculative process:




Memory Protection

In Virtual-Memory Based TLS

main process:



speculative process:




→ `mprotect(addr, 4096, PROT_READ|PROT_WRITE)`


Memory Protection

In Virtual-Memory Based TLS

main process:



speculative process:




→ `mprotect(addr, 4096, PROT_READ|PROT_WRITE)`


Memory Protection

In Virtual-Memory Based TLS

main process:



speculative process:



→ `mprotect(addr, 4096, PROT_READ|PROT_WRITE)`


TLS Implementations

S-TLS

VM-TLS


	user space library	instrume ntation	setup cost	execution overhead	commit cost	granula- rity	I/O isolation
S-TLS	✓	✓	\$	\$\$\$ (per access)	\$\$\$	✓ (arbitrary)	—
VM-TLS	✓	—	\$\$\$	\$\$ (per page)	\$\$	✗ (memory page)	—

TLS Implementations



	user space library	instrume ntation	setup cost	execution overhead	commit cost	granula- rity	I/O isolation
S-TLS	✓	✓	\$	\$\$\$ (per access)	\$\$\$	✓ (arbitrary)	—
VM-TLS	✓	—	\$\$\$	\$\$ (per page)	\$\$	✗ (memory page)	—

TLS Implementations



	user space library	instrume ntation	setup cost	execution overhead	commit cost	granula- rity	I/O isolation
S-TLS	✓	✓	\$	\$\$\$ (per access)	\$\$\$	✓ (arbitrary)	—
U-TLS	✓	—	\$\$\$	\$\$ (per page)	\$\$	✗ (memory page)	—
K-TLS	—	—	\$\$	\$ (per page)	\$	✗ (memory page)	✓

Memory Protection In Kernel Space

main process:

VMA 1: r-x


VMA 2: rw-



speculative process:

VMA 1: r-x

VMA 2: rw-



Memory Protection In Kernel Space

main process:

VMA 1: r-x


VMA 2: rw-



speculative process:

VMA 1: r-x

VMA 2: rw-



Memory Protection In Kernel Space

main process:

VMA 1: r-x


VMA 2: rw-



speculative process:


VMA 1: r-x

VMA 2: rw-





Evaluation

Spawn Time




Evaluation

Random Memory Writes




Evaluation

Strassen Matrix Multiplication




Evaluation

Strassen Matrix Multiplication





Evaluation

Strassen Matrix Multiplication




Evaluation

Strassen Matrix Multiplication



Evaluation

Cilk Suite



Summary

- K-TLS: Thread-Level Speculation with Kernel Support
 - superior to software-only approaches like STM
 - superior to virtual-memory based method in user space
 - full isolation even for system calls
 - no instrumentation → easy to integrate
 - open source: <https://github.com/hammacher/k-tls>